

Memor — Adversarial Product Review

Local shared-memory & token-compression layer for agentic coding tools · evaluated as a global context layer for Codex

Reviewer: Head of Product (evaluation lens)
Version: v0.14.0 · MIT · 1,705 tests
Date: 2026-09-07
Data: one developer store, 40,810 artifacts

Bottom line. Memor's differentiator is real and rare: it builds cross-tool memory from transcripts agents *already write to disk*, so a tool contributes memories without any plugin, MCP, or config — Goose wrote 1,256 memories having never heard of memor. Compression is honestly measured and net-of-cache. **But for the Codex-specific question the verdict is: not yet proven.** Codex is a heavy *writer* (4,602 memories) whose reading is *unmeasured* rather than broken: its 38 zero-hit recalls all predate the 2026-08-05 fix that taught the proxy to resolve a project, and only 2 recalls have been made since. The real gap is that **one agent does the reading** — Claude Code — while Cursor and jcode silently stopped weeks ago, and the product still cannot show it makes any agent measurably better at tasks.

VERIFIABLE STRENGTHS (MARKETABLE)

Zero-cooperation capture unique
Reads transcripts agents write for their own reasons. No save-tool call, no per-harness plugin, no screen capture.

Evidence: Goose 1,256 & Kimi 95 memories with no memor integration installed. Reproducible: artifacts-by-source.

Genuinely cross-tool, scoped by project
A decision in Claude Code is recalled in Cursor/Codex. Scope is project, not agent, so memory crosses tools by default.

Evidence: 40,810 artifacts across 33 projects; 6 writer agents in one store.

Compression is real and self-funding
Lifetime **2,453,627 tokens** saved; recall injected 1.15M. Compression nets positive after paying for memory.

Evidence: proxy_savings ledger (non-passthrough). Reproduce: memor compression-worth.

Honest, falsifiable measurement rare
Savings reported **net of prompt cache** (cache writes priced against reads); says *unmeasured* rather than assuming zero. Every headline has a replay command.
88.8% on a pytest -v payload; 47.7% on the Bash output it engages; all via memor hook-worth.

Local-first, no vendor lock-in
Runs entirely on localhost, no memor API key, forwards your existing provider credentials. Local embedding (~2ms).

Evidence: proxy forwards Anthropic/OpenAI creds; hybrid dense+BM25 retrieval with relevance gate.

Retriever works in isolation
95.0% any-hit, 86.7% all-gold on LongMemEval_S (n=120).
Reproduce: memor eval-longmemeval. See caveat — scores retrieval, not task success.

DRAWBACKS TO FIX BEFORE WE'D SHIP IT

Only one agent actually reads blocker
Claude Code made **689 of 727** recalls since 2026-08-05. Cursor last read 2026-08-11, jcode 2026-08-22, and nothing surfaced either. Codex's 38 zero-hit recalls are *pre-fix June* rows, not evidence of a broken integration — a lifetime total spanning a fix reported retired code as current. Cross-tool memory cannot be judged when one tool does the reading. *The one integration OpenAI cares about is the least proven.*

No task-outcome benchmark blocker
Cannot show it makes an agent *better*. Its own counterfactual win rate is **8.6%** (was reported 63.8% before the harness was fixed). SWE-bench can't be borrowed (no prior session history).
Design-only plan exists; benchmark unbuilt.

2 in 5 recalls return nothing
39.2% miss rate (1,758 / 4,485). Some are correct declines, but the share is high enough to erode trust.
recall_log, all agents. Claude 1,544/3,986 miss.

Quality graded on a razor-thin sample
Only 453 "used" verdicts vs 40,810 artifacts (~1%). Cross-tool — the whole point — is barely judged. Rejection detector was blind before Aug 23.
recall_outcomes: 453 used / 1,235 unused / 133 pending.

Silent reliability failures
Proxy savings recorded only **17 days total** and nothing since Aug 22, and two agents stopped reading without a word. A layer you forget is running is a layer that silently stops. *Partly addressed: memor doctor now names any agent that stopped reading.*
proxy_savings spans 2026-08-01 → 08-22 only.

Blended compression looks tiny (0.8%)
Arithmetically honest (tool results are 33% of a request) but a hard marketing story. Real ceiling ~3.9% of a whole request.
Only tool results are rewritable; args/convol/images are not.

Security posture is deny-list based
Credential *filenames* are refused and structured secrets redacted, but bare-digit secrets (2FA codes) once passed through. Filename heuristics are brittle.
Fixed reactively after a real backup-codes file leaked past redaction.

What would make this marketable for Codex: (1) ~~Land Codex recall via the Responses API~~ — **misdiagnosed**. Recall already resolves the project from the system prompt and tool paths (fixed 2026-08-05); Codex's MCP registration simply omitted its agent label, so its reads were filed under jcode. Both now fixed and tested. (2) Get more than one agent reading, then ship a task-outcome benchmark with prior session history; a credible win-rate is the only claim buyers act on, and it cannot be measured on a store only Claude Code reads. (3) **Done:** memor doctor reports per-agent staleness and withholds hit rates drawn from pre-fix data. (4) Grow graded cross-tool verdicts from ~1% to a defensible sample. (5) Move security from filename deny-list to content-based secret detection.

Correction (2026-09-07, after this review shipped). The original draft called Codex "write-only in practice" on the strength of 38 recalls with 0 hits. Every figure was pulled live and the conclusion was still wrong: all 38 rows date from June, before the fix that gave the proxy a project to search, so the metric described code that no longer runs. A lifetime aggregate hides the moment a bug was fixed. Any per-agent statistic here is now windowed against the last behaviour change, which is what memor doctor enforces.